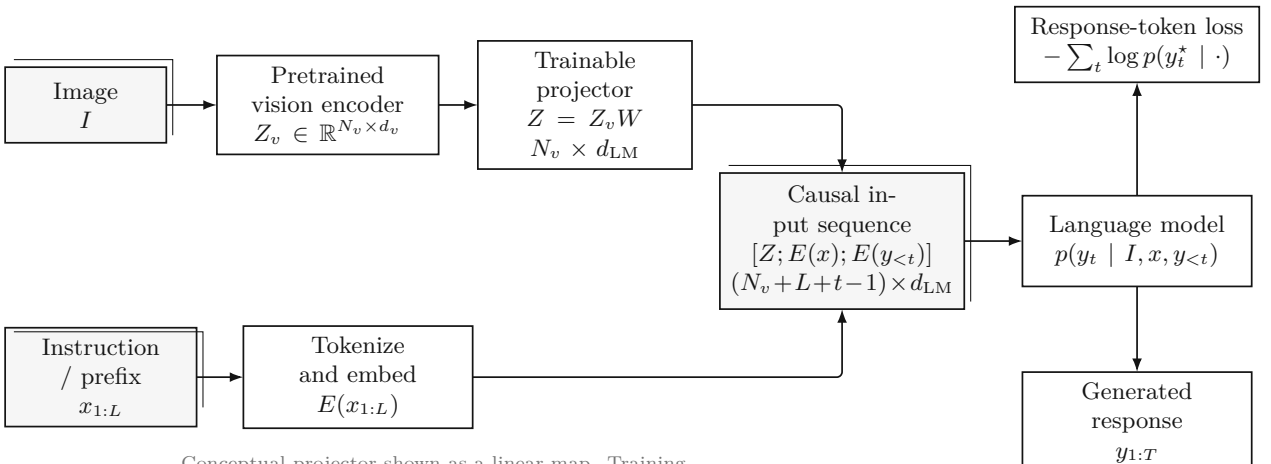
Vision-language projection bridge

LLaVA-style conceptual graph: project visual patch features into the language-model embedding space.



Conceptual projector shown as a linear map. Training stages determine which pretrained weights are updated.

Projected image tokens and text embeddings enter one causal sequence; response tokens carry the language objective.