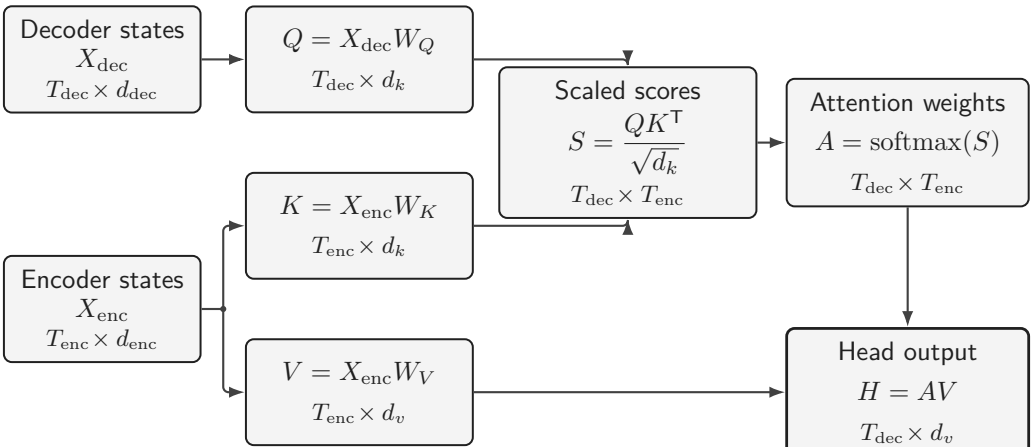
Single-head cross-attention

Decoder queries attend to encoder keys and values.



$$H = \text{softmax}\left(\frac{QK^{\mathsf{T}}}{\sqrt{d_k}}\right) V$$

Softmax normalizes each query row over encoder positions. Batch axis omitted.